






REVIEW: THE BAYESIAN LINEAR MODEL

Let $\mathcal{D} = \{(\mathbf{x}^{(1)}, y^{(1)}), \dots, (\mathbf{x}^{(n)}, y^{(n)})\}$ be a training set of i.i.d. observations from some unknown distribution.






Let $\mathbf{y} = (y^{(1)}, \dots, y^{(n)})^\top$ and $\mathbf{X} \in \mathbb{R}^{n \times p+1}$ be the design matrix where the i -th row contains vector $\mathbf{x}^{(i)}$.




REVIEW: THE BAYESIAN LINEAR MODEL



REVIEW: THE BAYESIAN LINEAR MODEL



REVIEW: THE BAYESIAN LINEAR MODEL



REVIEW: THE BAYESIAN LINEAR MODEL


Based on the posterior distribution

$$\boldsymbol{\theta} \mid \mathbf{X}, \mathbf{y} \sim \mathcal{N}(\sigma^{-2} \mathbf{A}^{-1} \mathbf{X}^{\top} \mathbf{y}, \mathbf{A}^{-1})$$



we can derive the predictive distribution for a new observation $\mathbf{x}_*$. The predictive distribution for the Bayesian linear model, i.e. the distribution of $\boldsymbol{\theta}^{\top} \mathbf{x}_*$, is

$$y_* \mid \mathbf{X}, \mathbf{y}, \mathbf{x}_* \sim \mathcal{N}(\sigma^{-2} \mathbf{y}^{\top} \mathbf{X} \mathbf{A}^{-1} \mathbf{x}_*, \mathbf{x}_*^{\top} \mathbf{A}^{-1} \mathbf{x}_*)$$

Please see the Deep Dive part for more details.





REVIEW: THE BAYESIAN LINEAR MODEL





For every test input $\mathbf{x}_*$, we get a distribution over the prediction y_* . In particular, we get a posterior mean (orange) and a posterior variance (grey region equals $\pm$ two times standard deviation).

REVIEW: THE BAYESIAN LINEAR MODEL



For every test input $\mathbf{x}_*$, we get a distribution over the prediction y_* . In particular, we get a posterior mean (orange) and a posterior variance (grey region equals $\pm$ two times standard deviation).

REVIEW: THE BAYESIAN LINEAR MODEL



For every test input $\mathbf{x}_*$, we get a distribution over the prediction y_* . In particular, we get a posterior mean (orange) and a posterior variance (grey region equals $\pm$ two times standard deviation).

SUMMARY: THE BAYESIAN LINEAR MODEL

- By switching to a Bayesian perspective, we do not only have point estimates for the parameter θ , but whole **distributions**
- From the posterior distribution of θ , we can derive a predictive distribution for $y_* = \theta^\top \mathbf{x}_*$.
- We can perform online updates: Whenever datapoints are observed, we can update the **posterior distribution** of θ

Next, we want to develop a theory for general shape functions, and not only for linear function.

