


COST CURVES

- Directly plot the misclassification costs / error (in terms of prior probs)
- Might be easier to interpret than ROC, especially in case of different misclassification costs or priors


Example:

- f_1 and f_2 with intersecting ROC curves
- f_2 dominates first, then f_1

BUT: Unclear for which thresholds, costs or class distribs f_2 better than f_1



Nathalie Japkowicz (2004): Evaluating Learning Algorithms : A Classification Perspective. (p. 125)




COST CURVES

Simplifying assumption: equal misclassif costs, i.e., $cost_{FN} = cost_{FP}$

⇒ Expected misclassif cost reduces to misclassif error rate

With law of total prob, we write error rate as function of π_+ :

$$\begin{aligned}\rho_{MCE}(\pi_+) &= (1 - \pi_+) \cdot \mathbb{P}(\hat{y} = 1|y = 0) + \pi_+ \cdot \mathbb{P}(\hat{y} = 0|y = 1) \\ &= (1 - \pi_+) \cdot FPR + \pi_+ \cdot FNR \\ &= (FNR - FPR) \cdot \pi_+ + FPR\end{aligned}$$



Confusion matrix

	True class	
	$y = 1$	$y = 0$
Pred. $\hat{y} = 1$	TP	FP
class $\hat{y} = 0$	FN	TN

Cost matrix




	True class	
	$y = 1$	$y = 0$
Pred. $\hat{y} = 1$	0	$cost_{FP}$
class $\hat{y} = 0$	$cost_{FN}$	0

COST CURVES

- Cost line of a classifier with slope ($FNR - FPR$) and intercept FPR :

$$\rho_{MCE}(\pi_+) = (FNR - FPR) \cdot \pi_+ + FPR$$

- Cost curves are point-line duals of ROC curves, i.e., a single classifier is represented by a point in the ROC space and by a line in cost space



Chris Drummond and Robert C. Holte (2006): Cost curves: An improved method for visualizing classifier performance. Machine Learning, 65, 95-130 ([URL](#)).



COST LINES

Cost line of a classifier with slope $(FNR - FPR)$ and intercept FPR :

$$\rho_{MCE}(\pi_+) = (FNR - FPR) \cdot \pi_+ + FPR$$



- Hard classifiers are points (TPR, FPR) in ROC space
- The cost line of a classifier connects (π_+, ρ_{MCE}) -points at $(0, FPR)$ and $(1, 1 - TPR)$
- Classifier 3 always dominates classifier 1
- Classifier 3 is better than classifier 2 when $\pi_+ < 0.7$

Cost lines plot different values of π_+ vs. $\rho_{MCE}(\pi_+)$



COST LINES - EXAMPLE

- Horizontal dashed line: worst classifier (100% error rate for all π_+)
 $\Rightarrow FNR = FPR = 1$
- x-axis: perfect classifier (0% error rate for all π_+) $\Rightarrow FNR = FPR = 0$





$$\rho_{MCE} = (FNR - FPR) \cdot \pi_+ + FPR$$

		Confusion matrix	
		True class	
Pred. class	$\hat{y} = 1$	TP	FP
	$\hat{y} = 0$	FN	TN

COST LINES - EXAMPLE

- Horizontal dashed line: worst classifier (100% error rate for all π_+)
 $\Rightarrow FNR = FPR = 1$
- x-axis: perfect classifier (0% error rate for all π_+) $\Rightarrow FNR = FPR = 0$
- Dashed diagonal lines: trivial classifiers, i.e., ascending diagonal always predicts negative instances ($\Rightarrow FNR = 1$ and $FPR = 0$) and vice versa





$$\rho_{MCE} = (FNR - FPR) \cdot \pi_+ + FPR$$

		Confusion matrix	
		True class	
Pred. class	$\hat{y} = 1$	$y = 1$	$y = 0$
	$\hat{y} = 0$	TP	FP
		FN	TN

COST LINES - EXAMPLE

- Horizontal dashed line: worst classifier (100% error rate for all π_+)
 $\Rightarrow FNR = FPR = 1$
- x-axis: perfect classifier (0% error rate for all π_+) $\Rightarrow FNR = FPR = 0$
- Dashed diagonal lines: trivial classifiers, i.e., ascending diagonal always predicts negative instances ($\Rightarrow FNR = 1$ and $FPR = 0$) and vice versa
- Descending/ascending bold lines: two families of classifiers *A* and *B* (represented by points in their respective ROC curves)






$$\rho_{MCE} = (FNR - FPR) \cdot \pi_+ + FPR$$

		Confusion matrix	
		True class	
Pred. class	$\hat{y} = 1$	$y = 1$	$y = 0$
	$\hat{y} = 0$	TP	FP
		FN	TN




VISUALIZE COST CURVE - LOWER ENVELOPE

- Left: ROC = TPR & FPR of a classifier for different prob thresholds
- Right: Corresponding cost lines
- Duality: For every ROC point we can construct the CC line, and vice versa.






VISUALIZE COST CURVE - LOWER ENVELOPE

- Left: ROC = TPR & FPR of a classifier for different prob thresholds
- Right: Corresponding cost lines
- Duality: For every ROC point we can construct the CC line, and vice versa.



VISUALIZE COST CURVE - LOWER ENVELOPE

- Left: ROC = TPR & FPR of a classifier for different prob thresholds
- Right: Corresponding cost lines
- Duality: For every ROC point we can construct the CC line, and vice versa.



VISUALIZE COST CURVE - LOWER ENVELOPE

- Left: ROC = TPR & FPR of a classifier for different prob thresholds
- Right: Corresponding cost lines
- Duality: For every ROC point we can construct the CC line, and vice versa.
- **Cost curve** (right, black) is lower envelope of **cost lines**
 $\triangleq$ pointwise minimum of error rate (as function of π_+)

